


Observed-effect reliability ceiling vs best model, per task

split-half pseudo-replicate reliability of the observed effect (Spearman-Brown full-sample); floor = best simple baseline



† T3 / T4: deposited best-model & floor are macro-averaged on a coarser grain (T3 dataset $n=5$; T4 modality-fold $n=2$) than the per-unit reliability (per KO-gene), so the per-unit ceiling can sit below the coarse-grain score.